

Intro to IT Security

CS306C—Fall 2022

Prof. Antonio R. Nicolosi

Antonio.Nicolosi@stevens.edu



Authentication & Authorization (Access Control)

Setup

- Let A, B denote two entities
- A will act as the **prover** (trying to prove her identity)
- B will act as the **verifier**

Generalities

- Suppose that A wants to authenticate herself to B
- This must be accomplished through some kernel of trust, based upon prior knowledge, *e.g.*,
 - a trusted public key
 - a shared secret

Authentication

The task of A proving to B that A is the owner of the corresponding private key or the shared secret, etc.

Ideally this proof should satisfy the following properties:


- Easy to create given the secret
- Hard to forge without the secret
- Remain hard to forge, no matter how many other proofs have been seen by an adversary

Authentication Via PRFs

Pseudo-Random Functions (PRFs) are a simple way to construct an authentication scheme.


Setup:

- $\mathcal{F} = \{f_i : \mathbb{Z}_n \rightarrow \mathbb{Z}_m \mid i \in \{0, 1\}^k\}$, a family of pseudorandom functions
- Shared secret: $s \in \{0, 1\}^k$
- Algorithm:



Authentication via Public-Key Cryptography

If Enc denotes the encryption algorithm for an asymmetric scheme and if h denotes a cryptographic hash function, we can provide authentication as follows:



Types of Authentication

- End user → End user (e.g., Alice & Bob)
- End user → Local computer (e.g., login)
- End user → Remote computer (e.g., web site login)
- Computer → Computer (e.g., DRM)
- Local computer → End user (e.g., prevent fake ATM)
- Remote computer → End user (e.g., anti-phishing)

Passwords

- Passwords are a classic way to authenticate
- Advantages:
 - Familiar; common
 - Easy to use
 - Easy to remember (sometimes...)

Problems with Passwords

- If password is sent in the clear, can be intercepted
- If password is encrypted, requires establishment of encryption key
- People choose bad passwords (e.g., not enough entropy)
- Passwords are easily observed (sniffed by spyware, or perhaps other humans)

Spy / Spyware Defenses

Use a randomized keypad; use mouse for input rather than keyboard



Storage of Plaintext Passwords

- If passwords are kept as plaintext, then a compromise of the system could be especially catastrophic
- Worse: Might be able to guess password in linear time!
 - Place first character of password as last byte on VM page
 - Make sure the next page is unmapped
 - Try all first letters, one by one
 - Got a page fault? Must have correct first letter
 - Repeat for second character of password, etc.

Hashed Passwords

One approach to mitigate the risks if the system is compromised is to hash the password first:

- The user computes and sends $h(pwd)$ for some one-way, collision-resistant hash h
- Operating system simply checks the hash

What could go wrong?

Attacks on Hashed Passwords

While better than storing passwords in the clear, this still leaves avenues for attacks.

Off-line Dictionary Attack:

- If h is known, one can compute $h(pwd)$ for a long list of common passwords
- If the file of hashed passwords is recovered, one can immediately check for matches and lookup the password

Example: Unix Hashed Passwords

- Store one-way function of password
 - Only hashes first 8 characters
 - Apply DES 25 times with password as key, starting from an all-zero input block
 - Key generally hard to recover from ciphertext
 - So put hashed password in world-readable /etc/passwd
 - To validate password, hash it and compare to stored hash
- Hash function “salted” with 12 extra bits for each user
 - Prevent attacker from building dictionary of hashes of common passwords
 - Prepend seed to hashed password for use in verification

Example: Unix Hashed Passwords

- Store one-way function of password
 - Only hashes first 8 characters
 - Apply DES 25 times with password as key, starting from an all-zero input block
 - Key generally hard to recover from ciphertext
 - So put hashed password in world-readable /etc/passwd
 - To validate password, hash it and compare to stored hash
- Hash function “salted” with 12 extra bits for each user
 - Prevent attacker from building dictionary of hashes of common passwords
 - Prepend seed to hashed password for use in verification

Unix Hashed Passwords: Weaknesses

- Off-line password guessing attacks
 - Attacker gets password file, guesses passwords at home
 - Computationally expensive, but undetectable by victims
- Fails to account for hardware & software improvements
 - DES crypt on 1977 VAX-11/780: 4 crypts/sec
 - Bitsliced implementation on 600 MHz alpha: 214,000 cr/sec
 - But users don't chose better passwords now than in 1977...
- Server still gets plaintext password at login
 - Attacker or server operator can modify login to record it
 - But users often use same password on different systems

Systems with passwords: Design Hints

- Any algorithm that uses passwords should:
 - Take a salt as input (prevent hash dictionaries)
 - Take a cost parameter as input
- Administrators need to increase hashing cost as hardware gets faster

Systems with passwords: Examples

- Bad examples: PGP & ssh private keys
 - Stored in user's home directory, encrypted with password
 - Given encrypted private key, can guess password off-line
 - Cost of guess even cheaper than password crypt!
- Good example: OpenBSD/FreeBSD bcrypt
 - Hashing password is exponential in cost parameter
 - Cost goes in hashed password along with salt

Network Passwords

- Many systems grant access through a password
- How to implement? Example:
 - Server stores user's password (or hash)
 - Client connects to server, sends username, password
 - Server compares password to stored version
 - Grants access if they match
- Is this a good approach?

Weaknesses

- How do you know you are talking only to the server?
 - Attacker might be eavesdropping on Ethernet
 - Attacker might mess with DNS (or IP routing) so you talk to wrong machine
- Eavesdropper will just read your password off the network
- Server knows your plaintext password
 - Or at least sees it at login time, even if stores hash
 - Extra-bad if people re-use passwords on multiple machines

s/key Password Authentication

- Goal: Protect against passive eavesdroppers
 - Also: Minimize harm of bad clients (e.g., public terminals)
- Idea: One time passwords, not valid after snooped
- Algorithm takes user's real password, p , a random "salt" s , and server machine name m
 - First one time password is $H^{(100)}(p, m, s)$ (for one way hash H)
 - Next password is $H^{(99)}(p, m, s)$, etc.
 - After 99 logins, must change salt or password
- Benefit: Very convenient
 - Carry list of one-time passwords or calculate on smartphone

Weaknesses

- Vulnerable to network attackers
 - Attacker impersonates server
 - Re-sends one-time password to real server
- Vulnerable to off-line password guessing
 - Attacker sees s and $H^{(n)}(p, m, s)$
 - Can verify guesses of p off-line—check against dictionary of common passwords
 - H not very expensive (should be $H^{(n)}(G(p, m, s))$ for expensive G —or perhaps increase n)
- Bad client can compromise session
 - Before logout, can insert command to create back door:
echo ssh-rsa AAAA...Fbg== badguy » .ssh/authorized_keys

Password-derived public keys

- Derive public key from user's password
 - E.g., use password as seed for pseudo-random generator to generate (public key, private key) pair
 - Client can regenerate private key given password
- Server stores public key for each user
- Use public key authentication, E.g.:
 - $S \rightarrow C : N_S$
 - $C \rightarrow S : K_u, \{C, S, K_S, u, session, N_S\}_{K_u^{-1}}$
 - K_S is server public key, used to authenticate server
 - S looks up public key K_u for credentials and checks signature

Weaknesses

- No salt
 - Users with same password will have same (public key, private key) pair
- No cost parameter
 - Can't take too long to log in
 - But over time generating key will get faster
- Public key is just like password hash
 - Eavesdropper will see key, can mount off-line attack
- No authentication of server to user
 - W/o accessing server, attacker can pretend server is giving bad answers

Just For Fun

Ever wondered how garage door openers work, or the remote control for the car door locks?

- PRF authentication would be nice, but it requires interaction...
- Building a transmitter into the car and a receiver into the remote seems like too much effort
- How to avoid it?

Non-Interactive Protocol

- Both the transmitter and receiver keep track of an index k
- Transmitter just sends PRF values in sequence: sends $a = f_s(k)$ and then increments k
- Transmitter verifies that $a = f_s(k)$: if so opens door and increments k

Problem: what happens if you hit the remote, but the receiver doesn't hear you?

Non-Interactive Protocol

To make the non-interactive protocol more usable, the designers sacrificed a little security:

- Both the transmitter and receiver keep track of an index k
- Transmitter again sends $a = f_s(k)$ and increments k
- Receiver checks to see if

$$a \in \{f_s(k + i)\}_{i=0}^{255}$$

- If so, then set $k = k + i$ where $a = f_s(k + i)$

Access Control

- Protected Entity: “object” O
- Active object: “subject” S
- An object can be
 - File
 - Directory of files
 - Executing program
 - Database entry
 - Data structure
 - System resource (printer, ...)

General Principles

- Fine-grained access control is desirable
 - *E.g.*, control access to files not just directories
 - More powerful, but of course more complex
- Least privilege
 - Grant minimum abilities necessary to complete task
- Closed vs.. open policies (a.k.a. default-on vs. default-off)
 - Closed: forbidden unless explicitly allowed
 - Open: allowed unless explicitly forbidden
- Conflict resolution
 - Prevent conflicts, or know how to deal with them

Access Control Policies

- Types of Policies
- Implementations

Access Control Policies

- Discretionary Access Control (DAC)
- Mandatory Access Control (MAC)
- Role-Based Access Control (RBAC)
- Not necessarily mutually exclusive
 - Can use different mechanisms for different resources
 - Can apply two policies, and allow access only if both allow

Discretionary Access Control (DAC)

- Control based on the identity of the user and access rules
- Rights can be delegated at user's *discretion*
- Most common (e.g., user, group, owner permissions on Unix)

Mandatory Access Control (MAC)

- Control over the policy is centralized
- Control based on comparing security labels with security clearances
- Delegation not allowed
- Originated in military applications
- Recently integrated into commodity systems (e.g., the “Mandatory Integrity Control” in Windows version > 6)

Role-Based Access Control (RBAC)

- Control based on user's (or program's) role (rather than their identity alone)
- User's right can change depending on their current role
- More recent proposal; adds flexibility and ease
- One could construct equivalent MAC structures, but RBAC is more naturally organized
- Examples include
 - Numerous database systems
 - FreeBSD
 - Microsoft Active Directory

DAC: Means of Implementation

- Access Control Matrix
- Access Control List
- Capabilities

Access Control Structure: An Analogy

- Think of the setup as a bipartite graph with labeled edges
 - The left vertexes are subjects, the right are objects and the edges (labeled) define the permissions
- Two common ways to represent graphs also apply to representing a control structure, and some of the same considerations apply
 - Matrix representation
 - List representation

Access Control Matrix

- Matrix indexed by all subjects and objects
 - Characterizes rights of each subject with respect to each object
- Formally: set of object O and subject S , set of possible rights
 - Objects: label columns of matrix
 - Subjects: label rows of matrix
 - Entry (s, o) contains the right for s to o
 - *E.g.*, read/write/execute

Access Control Matrix: Extensions

- Rights can be functions
 - “Actual” right depends on the system state
 - Equivalently, may depend on system history
- Rights can form hierarchies
 - *E.g.*, right X implies right Y
- How fine-grained the access control is depends on how fine-grained the rights are

Access Control Matrix: Drawbacks

- Number of subjects/objects is very large
- Matrix is usually sparse (most entries empty or contain some default value)
- One central matrix modified every time subjects/objects are created/deleted or rights are modified
- “Small” change can result in “many” changes to the access control matrix
 - *E.g.*, making a file publicly readable
- No support for dynamic access control
 - Does not describe if/how permissions are allowed to change

Access Control List (ACL)

- One access control list per object
- Can be viewed as storing the columns of the access control matrix with the appropriate object
- One list per object showing all subjects with access and their rights
- Possible to assign “default” rights to an object
 - Easy to make an object public
- Example: access based on user, group, and compartment
 - Use of wild cards
 - To specify a group of hosts, a 1 is used
 - To specify a full subnet with a wildcard: 172.16.30.0 0.0.0.255
 - This tells the router to match up the first three octets exactly, but the fourth octet can be any value

ACL in practice

- Full granularity may not be supported
 - *E.g.*, user-level, group-level, or only public/private?

Conflicts

- How to handle conflicts if two subjects give different permissions on an object
 - Disallow multiple owners
 - Allow access if any entry gives rights
 - Allow access only if no entry denies rights
 - Apply first applicable entry

Capabilities

- Can be viewed as storing the rows of the access control matrix with the appropriate subject
- Some burden for implementing protection placed on the user rather than just the OS
 - Analogy: user has a “ticket” which grants access to an object
 - A capability is an unforgeable token giving user access to an object and describing the level of allowable access
 - Capabilities can specify new types of rights

Capabilities: Two Approaches

- Ticket is held by OS, which returns to the subject a pointer to the ticket
- Ticket is held by the user, but protected from forgery by cryptographic mechanisms
 - How ...?
 - Two possibilities: ticket verified by the object or by the OS itself
 - Who holds the key in each case ...?

ACL vs. Capabilities

- Access control list
 - ACL associated with each object
 - Upon request, check user/group against the ACL
 - Relies on authentication of the user
- Capabilities
 - Can be passed from one user/process to another
 - Upon request, check validity of capability
 - No need to know the identity of the user/process making the request

ACL vs. Capabilities

- How would delegation be handled using ACLs vs. using capabilities?
 - ACL: run process using the name of the caller, or OS can explicitly support delegation
 - Capabilities: allow delegation “for free”
- How can a user make a file public when capabilities are used?
- How to revoke capabilities?

Capabilities: Advantages

- Better at enforcing “principle of least privilege”
 - Provide access to minimal resources, to the minimal set of subjects
- Allow much finer-grained control over subjects (process-level instead of user-level)

Capabilities: Disadvantages

- Overhead
- Revocation more difficult
- Controlling delegation more difficult
- Making files world-readable more difficult (impossible?)

MAC: Military Security Policy

- Primarily concerned with secrecy
- Objects given “classification” (rank; compartments)
- Subjects given “clearance” (rank; compartments)
- “Need to know” basis
 - Subject with clearance (r, C) dominates object with classification (r', C') only if $r \geq r'$ and $C' \subseteq C$
 - Classifications/clearance not necessarily hierarchical

Bell-LaPadula model

- Simple security condition: S can read O if and only if $l_o \leq l_s$
- *-property: S can write O if and only if $l_s \leq l_o$
 - Why?
- “Read down; write up”
 - Information flows upward

Dynamic Rights

- Could consider dynamic rights
 - Once a process reads a file at one security level, cannot write to any file at a lower security level

Basic Security Theorem

- If a system begins in a secure state, and always preserves the simple security condition and the *-property, then the system will always remain in a secure state
 - *I.e.*, information never flows down

Communicating Down

- How to communicate from a higher security level to a lower one?
- Dynamic rights
 - Max. security level vs. current security level
 - Maximum security level must always dominate the current security level
- Declassification
 - Reduce security level to write down
 - Security theorem no longer holds
 - Must rely on users to be security-conscious

Role-Based Access Control (**RBAC**)

- Access controls assigned based on *roles*
 - Can use an access matrix, where “subjects” are roles
- Users assigned to different roles
 - Can be static or dynamic
 - A user can have multiple roles assigned
 - Can use “access matrix” with users as rows, and roles as columns
 - Will, in general, be more compact than a full-blown access control matrix